

In order to discuss algorithms and compare them effectively, we need to start with a basic set of tools. Here, we explain these tools and provide a few examples. Rather than spend time honing our use of these tools, we will learn how to use them by applying them in our studies of actual algorithms. Lets start with a motivating example algorithm.

## Motivating Example: Insertion Sort

Before discussing any algorithm, it is good to write down a simple spec (as simple as possible) of the *problem* we are trying to use this algorithm to solve!! (Many descriptions of algorithms fail to do this and cause confusion. Let us not fall into this trap.)

The problem for today is the *sorting problem* whose input is an array of  $n$  numbers, and the output is an array of the given numbers sorted in increasing order. (Our notation for arrays is mostly standard. For example the notation  $A = [012572]$  asserts  $A$  is an array with six integers (indexed 0 to 5). The notation  $A[i]$  stands for the the  $i$ th element in the array. So  $A[0] = 0$ ,  $A[3] = 5$  and  $A[A[3]] = 2$ .)

The algorithm we wish to understand is the InsertionSort algorithm below.

---

**Algorithm 1** InsertionSort

---

```
Input: A
for  $i = 0$  to  $n - 2$  do
   $j = i$ 
  while  $j \geq 0$  and  $A[j + 1] < A[j]$  do
    swap  $A[j]$  and  $A[j + 1]$ 
     $j = j - 1$ 
  end while
end for
```

---

Given any algorithm we would want to prove it correctly solves the problem it was designed to solve. And then to analyze its speed or the “runtime”. While the focus of this lecture is mostly on the runtime analysis, let us first quickly deal with correctness.

To prove the correctness of the algorithm, we show by induction on  $i$  the claim that at the start of each iteration of the loop, the array elements  $A[0]..A[i]$  are sorted.

As a base case, when  $i = 0$ , there is only one element, and every one-element list is sorted.

For general  $i$ , we claim again by induction that at the beginning of the while loop, elements  $A[0]..A[i]$  with the exception of  $A[j+1]$  are sorted. and  $A[j+1] \leq A[j+2]$  if  $j+2 \leq i$ . Again this is true in the base case  $j = i-1$  since  $A[0]..A[i-1]$  were sorted at the end of the previous iteration of the outer loop. It is clear that this property is preserved by the conditional swap, yielding the induction step. Finally when break out of the loop, the inductive hypothesis combined with the condition  $i = -1$  or  $A[j] < A[j+1]$  imply  $A[0]..A[i]$  are sorted. And it remains after every step since  $A[j]$  as well as  $A[j+1]$  are smaller than all elements in  $A[j+2]..A[i]$  which are sorted, and so after the swap all of  $A[j]..A[i]$  are sorted. When we break out of the loop all of  $A[1]..A[i]$  are sorted, provind the inductive step for the outer induction.

When  $i = n-1$ , the statement we've proven by induction is that the whole list is in sorted order.

Now we turn to the question of how to analyze this algorithm. Specifically we can ask how many additions, comparisons, variable assignments, and swaps does InsertionSort take? This depends on how many times the inner loop runs, which may be as few as  $n$  times (if the input list is in sorted order) or as many as  $\binom{n}{2}$  times (if the input list is sorted in reverse). If we let that number of times be  $s$ , then there are

1. between  $2s+1$  and  $2s+n$  additions, depending on whether we recompute  $n-2$  each time or cache it,
2.  $n+2s$  comparisons
3.  $n+s$  assignments to variables
4.  $n+2s$  additions
5.  $s$  swaps.

Keeping track of each individual component exactly and then adding them up to compute the total number of steps an algorithm takes can get tediously detailed and messy. Most importantly it takes away the qualitative understanding of the running time that will help us compare different algorithms. This is why we work with the *asymptotic worst-case* running time of an algorithm. This will make the analysis of the the runtime of algorithms much easier. Specifically:

1. We consider only the *worst case* running time of an algorithm on inputs of length  $n$ . In InsertionSort, we can show that  $s \leq \binom{n}{2}$  and this can be achieved. This will allow us to work with just the one parameter  $n$ .
2. We measure the running time only up to constant factors. Comparisons, additions, swaps, and variable assignments may take different amounts of time, and may even take different relative amounts of time on different machines or in different languages, but each of them takes a constant amount of time.

We will introduce this notation more carefully next, but once done we will be able to make the simple statement “The running time of InsertionSort is  $O(n^2)$  (or  $\Theta(n^2)$ )”.)

## **$O$ (Big-Oh) Notation**

When measuring, for example, the number of steps an algorithm takes in the worst case, our result will generally be some function  $T(n)$  of the input size,  $n$ . One might imagine that this function may have some complex form, such as  $T(n) = 4n^2 - 3n \log n + n^{2/3} + \log^3 n - 4$ . In very rare cases, one might wish to have such an exact form for the running time, but in general, we are more interested in the rate of growth of  $T(n)$  rather than its exact form.

The  $O$  notation was developed with this in mind. With the  $O$  notation, only the fastest growing term is important, and constant factors may be ignored. More formally:

**Definition 3.1** We say for non-negative functions  $f(n)$  and  $g(n)$  that  $f(n)$  is  $O(g(n))$  if there exist positive constants  $c$  and  $N$  such that for all  $n \geq N$ ,

$$f(n) \leq cg(n).$$

Let us try some examples. We claim that  $2n^3 + 4n^2$  is  $O(n^3)$ . It suffices to show that  $2n^3 + 4n^2 \leq 6n^3$  for  $n \geq 1$ , by definition. But this is clearly true as  $4n^3 \geq 4n^2$  for  $n \geq 1$ . (**Exercise:** show that  $2n^3 + 4n^2$  is  $O(n^4)$ .)

We claim  $10 \log_2 n$  is  $O(\ln n)$ . This follows from the fact that  $10 \log_2 n \leq (10 \log_2 e) \ln n$ .

If  $T(n)$  is as above, then  $T(n)$  is  $O(n^2)$ . This is a bit harder to prove, because of all the extraneous terms. It is, however, easy to see;  $4n^2$  is clearly the fastest growing term, and we can remove the constant with  $O$  notation. Note, though, that  $T(n)$  is  $O(n^3)$  as well! The  $O$  notation is not tight, but more like a  $\leq$  comparison.

Similarly, there is notation for  $\geq$  and  $=$  comparisons.

**Definition 3.2** We say for non-negative functions  $f(n)$  and  $g(n)$  that  $f(n)$  is  $\Omega(g(n))$  if there exist positive constants  $c$  and  $N$  such that for all  $n \geq N$ ,

$$f(n) \geq cg(n).$$

We say that  $f(n)$  is  $\Theta(g(n))$  if both  $f(n)$  is  $O(g(n))$  and  $f(n)$  is  $\Omega(g(n))$ .

The  $O$  notation has several useful properties that are easy to prove.

**Lemma 3.3** If  $f_1(n)$  is  $O(g_1(n))$  and  $f_2(n)$  is  $O(g_2(n))$  then  $f_1(n) + f_2(n)$  is  $O(g_1(n) + g_2(n))$ .

**Proof:** There exist positive constants  $c_1, c_2, N_1$ , and  $N_2$  such that  $f_1(n) \leq c_1 g_1(n)$  for  $n \geq N_1$  and  $f_2(n) \leq c_2 g_2(n)$  for  $n \geq N_2$ . Hence  $f_1(n) + f_2(n) \leq \max\{c_1, c_2\}(g_1(n) + g_2(n))$  for  $n \geq \max\{N_1, N_2\}$ . ■

**Exercise:** Prove similar lemmata for  $f_1(n)f_2(n)$ . Prove the lemmata when  $O$  is replaced by  $\Omega$  or  $\Theta$ .

Finally, there is a bit for notation corresponding to  $\ll$ , when one function is (in some sense) much less than another.

**Definition 3.4** We say for non-negative functions  $f(n)$  and  $g(n)$  that  $f(n)$  is  $o(g(n))$  if

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0.$$

Also,  $f(n)$  is  $\omega(g(n))$  if  $g(n)$  is  $o(f(n))$ .

We emphasize that the  $O$  notation is a tool to help us analyze algorithms. It does not always accurately tell us how fast an algorithm will run in practice. For example, constant factors make a huge difference in practice (imagine increasing your bank account by a factor of 10), and they are ignored in the  $O$  notation. Like any other tool, the  $O$  notation is only useful if used properly and wisely. Use it as a guide, not as the last word, to judging an algorithm.

Returning briefly to InsertionSort, we can now say that the total number of steps that it takes is  $O(s + n)$  where  $s = O(\binom{n}{2})$ , which simplifies to a run time of  $O(n^2)$  for InsertionSort. We stress that this is an upper bound on the number of steps takes by the worst-case input. It is a good exercise to prove that there exists an input array  $A$  that takes  $\Omega(n^2)$  steps with InsertionSort. Another good exercise is to produce an input that takes only  $O(n)$  steps. PS1 gives a more detailed version of this exercise as an optional question.

We now turn to another common tool in analyzing algorithms, especially recursive ones.

## Recurrence Relations

A recurrence relation defines a function using an expression that includes the function itself. For example, the Fibonacci numbers are defined by:

$$F(n) = F(n-1) + F(n-2), \quad F(1) = F(2) = 1.$$

This function is well-defined, since we can compute a unique value of  $F(n)$  for every positive integer  $n$ .

Note that recurrence relations are similar in spirit to the idea of induction. The relations defines a function value  $F(n)$  in terms of the function values at smaller arguments (in this case,  $n-1$  and  $n-2$ ), effectively reducing the

problem of computing  $F(n)$  to that of computing  $F$  at smaller values. Base cases (the values of  $F(1)$  and  $F(2)$ ) need to be provided.

Finding exact solutions for recurrence relations is not an extremely difficult process; however, we will not focus on solution methods for them here. Often a natural thing to do is to try to guess a solution, and then prove it by induction. Alternatively, one can use a symbolic computation program (such as Maple or Mathematica); these programs can often generate solutions.

We will occasionally use recurrence relations to describe the running times of algorithms. For our purposes, we often do not need to have an exact solution for the running time, but merely an idea of its asymptotic rate of growth. For example, the relation

$$T(n) = 2T(n/2) + 2n, \quad T(1) = 1$$

has the exact solution (for  $n$  a power of 2) of  $T(n) = 2n \log_2 n + n$ . (**Exercise:** Prove this by induction.) But for our purposes, it is generally enough to know that the solution is  $\Theta(n \log n)$ .

The following theorem is extremely useful for such recurrence relations:

**Theorem 3.5** *The solution to the recurrence relation  $T(n) = aT(n/b) + cn^k$ , where  $a \geq 1$  and  $b \geq 2$  are integers and  $c$  and  $k$  are positive constants satisfies:*

$$T(n) \text{ is } \begin{cases} O(n^{\log_b a}) & \text{if } a > b^k \\ O(n^k \log n) & \text{if } a = b^k \\ O(n^k) & \text{if } a < b^k. \end{cases}$$

Finally, before returning to the sorting problem with a fresh algorithm, we introduce some basic data structures.

## Data Structures

We shall regard integers, real numbers, and bits, as well as more complicated objects such as arrays and lists, as primitive data structures. Recall that a list is just an ordered sequence of arbitrary elements.

$$\text{List } q := [x_0, x_1, \dots, x_{n-1}].$$

$x_0$  is called the head of the list.

$x_{n-1}$  is called the tail of the list.

$n = |q|$  is the size of the list.

We denote by  $\circ$  the concatenation operation. Thus  $q \circ r$  is the list that results from concatenating the list  $q$  with the list  $r$ .

The operations on lists that are especially important for our purposes are:

|                       |   |
|-----------------------|---|
| $\text{head}(q)$      | $\text{return}(x_0)$                                      |
| $\text{push}(q, x)$   | $q := [x] \circ q$  |
| $\text{pop}(q)$       | $q := [x_1, \dots, x_{n-1}], \text{return}(x_0)$          |
| $\text{inject}(q, x)$ | $q := q \circ [x]$  |
| $\text{eject}(q)$     | $q := [x_0, x_1, \dots, x_{n-2}], \text{return}(x_{n-1})$ |
| $\text{size}(q)$      | $\text{return}(n)$  |

The head, pop, and eject operations are not defined for empty lists. Appropriate return values (either an error, or an empty symbol) can be designed depending on the implementation.

A *stack* is a list that supports operations head, push, pop.

A *queue* is a list that supports operations head, inject and pop.

A *deque* supports all these operations.

Note that we can implement lists either by arrays or using pointers as the usual linked lists. Arrays are often faster in practice, but they are often more complicated to program (especially if there is no implicit limit on the number of items). In either case, each of the above operations can be implemented in a constant number of steps.

## Application: Mergesort

For the rest of the lecture, we will review the procedure mergesort. The input is a list of  $n$  numbers, and the output is a list of the given numbers sorted in increasing order. The main data structure used by the algorithm will be a queue. We will assume that each queue operation takes 1 step, and that each comparison (is  $x > y$ ?) takes 1 step. We will show that mergesort takes  $O(n \log n)$  steps to sort a sequence of  $n$  numbers.

The procedure mergesort relies on a function merge which takes as input two *sorted* (in increasing order) lists of numbers and outputs a single sorted list containing all the given numbers (with repetition).

```
function merge (s,t)
  list s,t
  if s = [ ] then return t
```

```

    else if  $t = []$  then return  $s$ 
    else if  $\text{head}(s) \leq \text{head}(t)$  then  $u := \text{pop}(s)$ 
        else  $u := \text{pop}(t)$ 
    return inject( $u$ , merge( $s, t$ ))
end merge

```

```

function mergesort ( $s$ )
    list  $s, q$ 
     $q = []$ 
    for  $x \in s$ 
        inject( $q, [x]$ )
    rof
    while  $\text{size}(q) \geq 2$ 
         $u := \text{pop}(q)$ 
         $v := \text{pop}(q)$ 
        inject( $q$ , merge( $u, v$ ))
    end
    if  $q = []$  return  $[]$ 
    else return  $q(0)$ 
end mergesort

```

The correctness of the function merge (i.e., that it “takes as input two *sorted* (in increasing order) lists of numbers and outputs a single sorted list containing all the given numbers (with repetition)”) is proved by induction using the following fact: the smallest number in the input is either  $\text{head}(s)$  or  $\text{head}(t)$ , and must be the first number in the output list. Suppose  $\text{head}(s)$  is the smallest element. Then this element should be the first element of the output list. Further, the rest of the output list is just the list obtained by merging  $s$  (after popping the first element) and  $t$ . Finally, merge works correctly by induction on this new pair of inputs (food for thought: what is the induction variable?) since  $t$  is sorted and so is  $s$  after popping the first element. (The case that  $\text{head}(t)$  is smaller is similar and omitted.)

The number of steps for each invocation of function merge is  $O(1)$  steps. Since each recursive invocation of merge removes an element from either  $s$  or  $t$ , it follows that function merge halts in  $O(|s| + |t|)$  steps.

**Question:** Can you design an iterative (rather than recursive) version of merge? How much time does it take? Which version would be faster in practice— the recursive or the iterative?

The iterative algorithm mergesort uses  $q$  as a queue of lists. (Note that it is perfectly acceptable to have lists of lists (or a queue of lists)!) It repeatedly merges together the two lists at the front of the queue, and puts the resulting list at the tail of the queue.

$$Q : [[7, 9], [1, 4], [6, 16], [2, 10] * [3, 11, 12, 14], [5, 8, 13, 15]]$$

$$Q : [[6, 16], [2, 10] * [3, 11, 12, 14], [5, 8, 13, 15], [1, 4, 7, 9]]$$

Figure 3.1: One step of the mergesort algorithm.

To prove the correctness of the algorithm, we claim that after initialization, every list in the queue  $q$  is a sorted list and the multiset of elements in the union of these lists is the input list  $s$ .

This correctness claim easily from the fact that we start with sorted lists of length 1 each, and merge them in pairs to get longer and longer sorted lists, until only one list remains. Furthermore, the elements in the list  $u$  and  $v$  were popped from  $q$  and then reinserted after merging thus preserving the multiset of elements.

Note that the correctness claim does not rely on the input set having  $2^k$  for integer  $k$ . However we recommend that to understand the algorithm for the first time, it is better to only consider input lengths that are a power of two. In fact, in what follows we will continue to assume that  $s$  has  $2^k$  elements, to make the runtime analysis cleaner.

To analyze the running time of this algorithm, let us place a special marker  $*$  initially at the end of the  $q$ . Whenever the marker  $*$  reaches the front of  $q$ , and is either the first or the second element of  $q$ , we move it back to the end of  $q$ . Thus the presence of the marker  $*$  makes no difference to the actual execution of the algorithm. Its only purpose is to partition the execution of the algorithm into phases: where a phase is the time between two successive visits of the marker  $*$  to the end of the  $q$ . Then we claim that the total time per phase is  $O(n)$ . This is because each phase just consists of pairwise merges of disjoint lists in the queue. Each such merge takes time proportional to the sum of the lengths of the lists, and the sum of the lengths of all the lists in  $q$  is  $n$ . On the other hand, the number of lists is halved in each phase, and therefore the number of phases is at most  $\log n$ . Therefore the total running time of mergesort is  $O(n \log n)$ .

An alternative analysis of mergesort depends on a recursive, rather than iterative, description. Suppose we have an operation that takes a list and splits it into two equal-size parts. (We will assume our list size is a power of 2, so that all sublists we ever obtain have even size or are of length 1.) Then a recursive version of mergesort would do the following:

```
function mergesort ( $s$ )
  list  $s, s_1, s_2$ 
  if size( $s$ ) = 1 then return( $s$ )
  split( $s, s_1, s_2$ )
   $s_1$  = mergesort( $s_1$ )
```



```

    s2 = mergesort(s2)
    return(merge(s1, s2))
end mergesort

```

Here `split` splits the list  $s$  into two parts of equal length  $s_1$  and  $s_2$ . The correctness follows easily from induction.

Let  $T(n)$  be the number of comparisons mergesort performs on lists of length  $n$ . Then  $T(n)$  satisfies the recurrence relation  $T(n) \leq 2T(n/2) + n - 1$ . This follows from the fact that to sort lists of length  $n$  we sort two sublists of length  $n/2$  and then merge them using (at most)  $n - 1$  comparisons. Using our general theorem on solutions of recurrence relations, we find that  $T(n) = O(n \log n)$ .

**Question:** The iterative version of mergesort uses a queue. Implicitly, the recursive version is using a stack. Explain the implicit stack in the recursive version of mergesort.

**Question:** Solve the recurrence relation  $T(n) = 2T(n/2) + n - 1$  exactly to obtain an upper bound on the number of comparisons performed by the recursive mergesort variation.